

An AI agent that asks before it acts

An agent that works the inbound lead inbox for Studio Lind, a design agency invented for this demo.

The problem

Leads arrive through a form and sit for days. Automating replies is risky: an agent can answer spam, misjudge a big lead, or send something nobody approved, and afterwards no one can explain why.

What I built

- Eight tools per lead: fetch, enrich, score, guardrails, draft, approval, CRM row, team summary
- Leads scoring 65 or more, or flagged by a guardrail, wait for a person to approve
- Guardrails block spam before the model is called
- Deterministic score where every point carries a written reason, shown in the drawer
- Every tool call kept with the lead and downloadable as JSON, failed steps included

Measured results

40 calls


tool calls in a full run over 7 seeded leads, every one logged

18 tests

pass with no API key, including reasons always summing to the score

2 bugs

found by my own checks: a budget like "about 40k" was discarded, and a form field named action shadowed form.action



The lead list, scored, sourced and waiting

The draft and the agent's working

Spam blocked before a reply

The replies in these screenshots were written by claude-haiku-4-5 through the local Claude CLI, which pays a process startup cost of around 30 seconds per lead. On the hosted API the same draft comes back in a few seconds. With no key at all the agent falls back to a template, which is how the tests run.

Stack

Python 3

FastAPI

Jinja2

SQLite

Anthropic API

Code github.com/vanthienha199/inbound-lead-agent